

Data card

source: outputs/reports/data_card_v1_20250815.md

Data Card — Synthetic Demo Dataset

- Version: 1.0 (2025-08-15)
- Generator: v0.1.0, seed 42, deterministic
- Mode: DEMO (synthetic only)

Sample formation

- 1000 billing lineages simulated across three fictional site contracts (CPL/TRI/NDX) with distinct ownership profiles.
- Business window 2022-07-01T00:00:00+00:00 to 2023-11-15T00:00:00+00:00; per-site local timezones with DST transitions; weekday-dominant rhythms; same-day entry ~45%, lagged entry lognormal (~2 days median), occasional very late arrivals.
- Outcome mix (ground truth): payment ~84%, still-open ~10%, no-payment close ~6%.
- Payer mix weighted (commercial ~55%, government ~32%, self ~8%, other ~5%).

Tables

| Table | Rows |

|---|---|

| bills | 1167 |

| events | 2145 |

| allocations | 883 |

| observation_windows | 6 |

Injected dirty cases (intercepted by validator)

- duplicate event imports, cross-org duplicate bill numbers, missing correction references, timezone-less timestamps, illegal (negative) amounts, non-USD currency.

Curation

- Invalid rows isolated, never deleted; original rows retained in raw layer.
- Hidden ground truth (true lineage membership, future payment times) stored under data/_ground_truth/, structurally excluded from model inputs.

Limitations

- Generator parameters are engineering loads, not claims about real distributions.
- No real patients, providers or payers are referenced.